

Question 2

(a) What is the probability that the first bootstrap observation is *not* the j-th observation from the original sample?

In bootstrap sampling, each observation is selected **with replacement** from the original sample of size n .

For any single draw (like the first bootstrap observation), the probability of selecting **any specific** observation (like the j-th) is $1/n$.

So, the probability at the **first** draw is **not** the j-th observation is:

$$1 - \frac{1}{n}$$

(b) What is the probability that the second bootstrap observation is not the j-th observation?

Just like part (a), each draw is independent and comes from the same distribution. So the probability is again:

$$1 - \frac{1}{n}$$

(c) What is the probability that the j-th observation is *not in the bootstrap sample*?

A bootstrap sample consists of n draws **with replacement**.

So the j-th observation is **not selected in any of the n draws** if it is not selected in draw 1, and not in draw 2, ..., and not in draw n .

Each draw independently avoids j-th observation with probability $1 - \frac{1}{n}$

, so the probability that **none of the n draws** select the j-th observation is:

$$\left(1 - \frac{1}{n}\right)^n$$

(d) When $n=5$, what is the probability that the j-th observation *is* in the bootstrap sample?

We already found that the probability the j-th observation is **not** in the bootstrap sample is:

$$\left(1 - \frac{1}{5}\right)^5 = \left(\frac{4}{5}\right)^5 = 0.32768$$

So the probability it **is** included at least once is:

$$1 - 0.32768 = 0.67232$$

(e) When $n=100$, what is the probability that the j -th observation is in the bootstrap sample?

$$\left(1 - \frac{1}{100}\right)^{100} \approx e^{-1} \approx 0.3679$$

So the probability it **is** included:


$$1 - 0.3679 \approx 0.6321$$


(f)

As $n \rightarrow \infty$, the probability that a given observation **is included** in the bootstrap sample approaches:

$$1 - \frac{1}{e} \approx 0.6321$$

(g) Plot the probability that the j -th observation is in the bootstrap sample for $n=1$ to 100000





(h)

0.632, confirming the theoretical result.

Question 3

(a)

In **k-fold cross-validation**, the data is randomly divided into **k approximately equal-sized parts**, called *folds*. The steps are:

1. **Split** the dataset into k folds.
2. For each fold:
 - Use that fold as the **validation set**.
 - Use the remaining $k-1$ folds as the **training set**.
 - **Train** the model on the training set.
 - **Evaluate** it on the validation set and record the error.
3. **Repeat** this process k times, each time using a different fold as the validation set.
4. **Average the k error estimates** to get the overall cross-validation estimate of test error.

This gives a more stable and reliable estimate of test performance than a single train/test split.

(b)

i. The validation set approach

Advantages:

- **Less variance:** Instead of relying on a single split, k-fold CV averages over multiple splits.
- **More efficient use of data:** Every observation is used for both training and validation (exactly once for validation and $k-1$ times for training).
- **More reliable test error estimate:** Especially useful when the dataset is small.

Disadvantages:

- **More computationally expensive:** Instead of training one model (as in the validation set approach), you train k models.

ii. Leave-One-Out Cross-Validation (LOOCV)

Advantages:

- **Computationally faster:** k-fold CV (with $k=5$ or 10) requires fewer model fits than LOOCV (which fits the model n times for n observations).
- **Lower variance** of test error estimate: LOOCV can have high variance because each training set is very similar (just one observation different).

Disadvantages:

- **Slightly more biased** than LOOCV: Because each training set in k-fold CV is smaller than the full dataset (LOOCV uses $n-1$ examples each time).
- **Less deterministic:** LOOCV gives a fixed result for a given dataset, while k-fold CV depends on the random partition (unless repeated multiple times).

Question 4

To estimate the **standard deviation** (i.e., uncertainty) of our prediction for a specific X , we can use **resampling methods**, especially the **bootstrap**.

Bootstrap Approach to Estimating Prediction Standard Deviation:

1. **Generate many bootstrap samples** from the original dataset:
 - Randomly sample (with replacement) from the dataset to create multiple new datasets of the same size.

2. **Fit the model** on each bootstrap sample.
3. **Make a prediction** for the same value of X using each fitted model.
4. **Collect all the predicted values** for X across the bootstrap samples.
5. **Calculate the standard deviation** of those predictions

Why this works:

The bootstrap mimics the process of sampling from the population. By observing how the prediction varies across different plausible datasets, we can estimate the **variability** (i.e., standard deviation) of the prediction due to sampling variation.

Notes:

- This standard deviation reflects **uncertainty due to sampling variability**, not model bias.
- It **does not account for model misspecification** — it assumes your model form is correct.
- For linear models, you can sometimes compute the prediction variance analytically, but bootstrap works for **any complex model**, including random forests, neural networks, etc.